

Klaudbiusz: Agent-Agnostic Tooling for Databricks Application Generation with Trajectory-Based Optimization and Composite Evaluation

Anonymous Author(s)

Abstract

We present Klaudbiusz, an open-source, agent-agnostic toolset for autonomous generation of Databricks data applications. Rather than building a specific agent, we provide reusable infrastructure—environment scaffolding, templates, and MCP tool integrations—that any agentic coding system can leverage. To continuously improve this toolset, we introduce a trajectory analyzer that processes agent execution traces via map-reduce LLM analysis, identifying friction points and generating actionable recommendations. To measure outcomes, we develop AppEval-100, a composite evaluation score built on four pillars: Reliability, SQL Quality (inspired by BIRD/Spider), Web Quality (inspired by WebArena), and Agentic DevX. Unlike existing benchmarks that evaluate SQL or web tasks in isolation, AppEval-100 provides the first composite evaluation for full-stack data applications, mapped to industry-standard DORA delivery metrics. We design a 100-prompt benchmark with $\pm 9\%$ confidence intervals across three difficulty tiers. Evaluated on 20 Databricks applications, we achieve 100% build/runtime success with 6–9 minute generation latency at \$0.74 per application. We release the complete framework including scaffolding tools, trajectory analyzer, evaluation harness, and MLflow integration.

CCS Concepts

- Software and its engineering → Software testing and debugging;
- Computing methodologies → Natural language processing.

Keywords

agent-agnostic tooling, agentic code generation, evaluation framework, trajectory optimization, Databricks

ACM Reference Format:

Anonymous Author(s). 2026. Klaudbiusz: Agent-Agnostic Tooling for Databricks Application Generation with Trajectory-Based Optimization and Composite Evaluation. In *Proceedings of the ACM Conference on AI and Agentic Systems (CAIS '26)*, May 26–29, 2026, San Jose, CA, USA. ACM, New York, NY, USA, 6 pages. <https://doi.org/XXXXXXX.XXXXXXX>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

CAIS '26, San Jose, CA, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM. ACM ISBN 978-1-4503-XXXX-X/26/05
<https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

The emergence of agentic coding systems—AI agents capable of autonomously generating, testing, and deploying software applications—represents a paradigm shift in software engineering. While benchmarks like HumanEval [2] and SWE-bench [5] have advanced our understanding of code generation quality, they focus primarily on functional correctness rather than production readiness.

The Infrastructure Gap. We observe that agent performance depends heavily on the quality of surrounding infrastructure: environment scaffolding, templates, tool integrations, and validation pipelines. Yet most work focuses on improving agents themselves rather than the reusable tooling that enables them. We argue for an *agent-agnostic* approach: build excellent infrastructure that any agent can leverage.

Core Principle. Our work is guided by a simple axiom: *If an AI agent cannot autonomously deploy code using provided tooling, the tooling needs improvement—not necessarily the agent.*

Three-Pillar Approach. We address this through three interconnected contributions:

- (1) **Agent-Agnostic Toolset.** We introduce Klaudbiusz, an open-source infrastructure for Databricks application generation comprising environment scaffolding, TypeScript/tRPC templates, MCP tool integrations, and containerized execution via Dagger. Any agentic system can leverage these tools.
- (2) **Trajectory Analyzer.** To continuously improve this toolset, we present a map-reduce LLM approach for analyzing agent execution traces. The analyzer identifies friction points, inefficient patterns, and tool failures, generating actionable recommendations for infrastructure improvement.
- (3) **Composite Evaluation (AppEval-100).** To measure outcomes, we develop a 4-pillar evaluation framework combining Reliability, SQL Quality, Web Quality, and Agentic DevX into a single composite score mapped to DORA delivery metrics.

Empirical Validation. Evaluated on 20 Databricks applications, we achieve 100% build/runtime success with 6–9 minute generation latency at \$0.74 per application.

2 Related Work

We survey evaluation approaches across four categories: function-level benchmarks, repository-level agent benchmarks, interactive agent environments, and evaluation harnesses. Table 1 summarizes key frameworks and identifies the gap our work addresses.

117 2.1 Function-Level Code Generation 118 Benchmarks

120 **HumanEval** [2] introduced pass@k evaluation on 164 Python functions,
121 becoming the standard metric for code generation. **MBPP** [1]
122 expanded to 974 problems but remains algorithm-focused. Both
123 benchmarks are now saturated—top models achieve >90% on HumanEval—
124 raising concerns about contamination and real-world relevance.

125 **BigCodeBench**¹ addresses these limitations with diverse library
126 calls and complex instructions, while **LiveCodeBench**² provides
127 dynamic, contamination-resistant evaluation. Recent work on **Hu-**
128 **manEval Pro and MBPP Pro** [16] introduces self-invoking code
129 generation to test progressive reasoning.

130 *Gap:* Function-level benchmarks do not evaluate deployment,
131 integration, or operational readiness.

132 2.2 Repository-Level Agent Benchmarks

134 **SWE-bench** [5] evaluates agents on 2,294 real GitHub issues, with
135 **SWE-bench Verified**³ providing 500 human-verified samples where
136 top models achieve ~72%. The harder **SWE-bench Pro**⁴ reveals
137 performance drops to ~23% on production-grade issues.

138 **SWE-agent** [14] demonstrates that custom agent-computer in-
139 terfaces significantly enhance performance through Docker-based
140 harnesses and tool augmentation. **REPOCOD**⁵ tests complete method
141 implementation with only 6% resolution rate, while **DevQuali-**
142 **tyEval**⁶ evaluates multi-language software engineering tasks on
143 private datasets to avoid contamination.

144 *Gap:* Repository-level benchmarks focus on patch correctness,
145 not autonomous deployment capability.

147 2.3 Text-to-SQL Benchmarks

149 For data-centric applications, text-to-SQL benchmarks provide crit-
150 ical evaluation capabilities. **Spider** [15] introduced cross-domain
151 evaluation with 10,181 questions across 200 databases, establishing
152 difficulty tiers (easy/medium/hard/extrahard). **BIRD** [9] scales to
153 12,751 question-SQL pairs across 95 databases (33.4GB), introducing
154 the *Valid Efficiency Score (VES)* that rewards both correctness and
query efficiency.

155 **Spider 2.0** [8] targets enterprise workflows with 632 real-world
156 tasks involving BigQuery and Snowflake, where even v1-preview
157 achieves only 21.3% (vs 91.2% on Spider 1.0). This dramatic per-
158 formance gap highlights the challenge of production-grade SQL
159 generation.

160 *Gap:* Text-to-SQL benchmarks evaluate query correctness in iso-
161 lation, not within deployed applications with UI, APIs, and DevOps
162 requirements.

164 2.4 Data Analytics Agent Benchmarks

166 Recent benchmarks evaluate agents on end-to-end data analysis.
167 **Tapilot-Crossing** [10] provides 1,024 human-machine interactions

169 ¹<https://github.com/bigcode-project/bigcodebench>

170 ²<https://livecodebench.github.io/>

171 ³<https://openai.com/index/introducing-swe-bench-verified/>

172 ⁴<https://scale.com/blog/swe-bench-pro>

173 ⁵<https://arxiv.org/abs/2410.21647>

174 ⁶<https://github.com/symflower/eval-dev-quality>

175 for interactive data analysis, testing multi-turn reasoning and visu-
176 alization. **InfiAgent-DABench** [4] offers 311 questions across 55
177 datasets for data analysis scenarios.

178 **InsightBench** [13] evaluates 100 business analytics tasks re-
179 quiring insight generation—the closest to our target domain. **DS-**
180 **1000** [7] benchmarks data science code generation across NumPy,
Pandas, and other libraries with 1,000 problems.

182 *Gap:* Data analytics benchmarks focus on insight generation
183 or code correctness, not full-stack application deployment with
184 Databricks integration.

186 2.5 Interactive Agent Benchmarks

187 **WebArena** [17] evaluates 812 web tasks across e-commerce, fo-
188 rumms, and content management, where best agents achieve 61.7%
versus 78% human performance. **GAIA** [12] tests general AI assis-
189 tants on 466 multi-step reasoning questions requiring tool use, with
agents reaching 80.7% versus 92% human baseline.

190 **AgentBench** [11] spans 8 environments (OS, databases, web) re-
191 vealing significant gaps between commercial and open-source mod-
192 els. **OSWorld**⁷ (NeurIPS 2024) benchmarks multimodal agents in
real computer environments, where recent advances have achieved
superhuman performance (76% vs 72% human baseline).

193 *Gap:* Interactive benchmarks evaluate general agent capabilities,
not software deployment pipelines.

201 2.6 Agent Evaluation Harnesses and 202 Frameworks

203 **Inspect AI**⁸ from the UK AI Safety Institute provides 100+ pre-
204 built evaluations with sandboxing, MCP tool support, and multi-
205 agent primitives. It has been adopted by frontier labs and safety
organizations for standardized agent evaluation.

206 **DeepEval**⁹ offers CI/CD integration with LLM-as-judge metrics
207 including task completion, tool correctness, and hallucination de-
208 tection. **Databricks Agent Evaluation** integrates with MLflow
209 for tracking groundedness, correctness, and coherence of agentic
210 applications.

211 The emerging **AgentOps** paradigm extends DevOps principles
212 to AI agents, addressing observability, tracing, and lifecycle man-
213 agement specific to autonomous systems.

217 2.7 DevOps and Deployment Metrics

218 **DORA metrics** [3]—deployment frequency, lead time, change fail-
219 ure rate, and mean time to restore—provide industry-standard mea-
220 sures of software delivery performance. **MLOps 2.0** architectures
221 integrate CI/CD with Continuous Data Validation (CDV) for reliable
ML delivery.

222 *Gap:* No existing framework combines code generation evalua-
223 tion with DORA-mapped deployment metrics and agentic DevX
224 scores (runability, deployability).

227 ⁷<https://os-world.github.io/>

228 ⁸<https://inspect.aisi.org.uk/>

229 ⁹<https://github.com/confident-ai/deepeval>

Table 1: Comparison of agent evaluation approaches. Klaudbiusz uniquely combines deployment-centric metrics with DORA mapping and trajectory-based optimization.

Framework	Code Gen	Deploy	DORA	DevX	Trajectory
HumanEval/MBPP	✓	–	–	–	–
SWE-bench	✓	–	–	–	–
WebArena/GAIA	–	–	–	–	–
Inspect AI	✓	–	–	–	–
DeepEval	✓	–	–	–	–
Klaudbiusz (Ours)	✓	✓	✓	✓	✓

3 The Klaudbiusz Toolset

3.1 Agent-Agnostic Design

Klaudbiusz provides reusable infrastructure that any agentic coding system can leverage, rather than a specific agent implementation. The toolset comprises:

- **Environment Scaffolding.** Pre-configured TypeScript + tRPC project templates with Databricks SDK integration, ensuring consistent structure across generated applications.
- **MCP Tool Integrations.** Model Context Protocol tools for file operations, database queries, and deployment actions that agents can invoke.
- **Containerized Execution.** Dagger-based sandboxed builds providing isolation and reproducibility across environments.
- **Validation Pipelines.** Automated checks for build, runtime, type safety, and deployment readiness.

3.2 Evaluation Design Principles

Our evaluation framework is built on two core principles:

Zero-Bias Metrics. All metrics are objective, reproducible, and automatable. We explicitly exclude subjective assessments of code quality, maintainability, or aesthetics.

Tooling-Centric Feedback. When agents fail, we ask “what tooling improvement would help?” rather than “what’s wrong with the agent?” This framing drives continuous infrastructure improvement.

3.3 The 13-Metric Rubric

We organize our metrics into four categories spanning core functionality, platform integration, agentic DevX, and generation efficiency.

3.3.1 Core Functionality (L1–L4, Binary).

- **L1: Build Success.** Project compiles; docker build exits with code 0.
- **L2: Runtime Success.** App starts and serves content; health check responds within 30s.
- **L3: Type Safety.** npx tsc –noEmit passes with zero errors.
- **L4: Tests Pass.** Unit/integration tests pass with coverage $\geq 70\%$.

3.3.2 Platform Integration (L5–L7, Binary).

- **L5: DB Connectivity.** Databricks connection works; queries execute without errors.

- **L6: Data Operations.** CRUD operations return correct data from tRPC procedures.
- **L7: UI Validation.** Frontend renders without errors (VLM verification).

3.3.3 Agentic DevX (D8–D9, 0–5 Score).

- **D8: Runability.** Can a sample AI agent run generated apps locally?
 - 0: install/start fails; missing scripts/env
 - 1–2: starts with manual tweaks
 - 3: starts cleanly with .env.example + documented steps
 - 4: starts with seeds/migrations via scripts
 - 5: + healthcheck endpoint + smoke test succeeds
- **D9: Deployability.** Can a sample AI agent deploy a generated app?
 - 0: no/broken Dockerfile
 - 1–2: image builds; container fails or healthcheck fails
 - 3: healthcheck OK; smoke 2xx
 - 4: + logs/metrics hooks present
 - 5: + automated rollback to prior known-good tag

3.3.4 Efficiency Metrics (E10–E13, Numeric).

- **E10: Tokens Used.** Total tokens (prompt + completion) for generation.
- **E11: Generation Time.** Time spent generating application (seconds).
- **E12: Agent Turns.** Number of conversation turns during generation.
- **E13: LOC.** Lines of code in generated application.

3.3.5 SQL Quality Pillar (S1–S4). Inspired by BIRD [9] and Spider [15], we evaluate SQL quality within generated applications:

- **S1: Execution Correctness (EX).** Fraction of generated SQL queries that execute without error and return expected results. Range: [0, 1].
- **S2: Valid Efficiency Score (VES).** Adapted from BIRD, rewards both correctness and query efficiency relative to a reference solution. Range: [0, 1].
- **S3: Query Complexity.** Distribution across difficulty tiers (easy/medium/hard/extrahard) based on Spider schema complexity.
- **S4: SQL Safety.** Absence of destructive operations (DROP, TRUNCATE), proper parameterization, and injection resistance. Range: [0, 1].

3.3.6 Web Quality Pillar (W1–W4). Inspired by WebArena [17] and VisualWebArena [6], we evaluate web/UI quality:

- **W1: Task Completion Rate.** Fraction of user-facing tasks that complete successfully (navigation, form submission, data display). Range: [0, 1].
- **W2: Visual Rendering Correctness.** VLM-assessed visual fidelity—charts render correctly, layouts are responsive, no broken elements. Range: [0, 1].
- **W3: Interactive Element Functionality.** Buttons, filters, and controls respond correctly to user input. Range: [0, 1].
- **W4: Accessibility Score.** WCAG compliance via automated tools (axe-core); keyboard navigation, ARIA labels, contrast ratios. Range: [0, 1].

3.4 AppEval-100 Composite Score

To enable automatic, comparable measurement across prompts and runs, we introduce **AppEval-100**—a single numeric index representing normalized readiness and agentic operability on a 0–100 scale. Unlike existing benchmarks that evaluate SQL correctness (BIRD/Spider) or web task completion (WebArena) in isolation, AppEval-100 provides the first composite evaluation combining all four pillars for full-stack data applications.

Step 1: Reliability Pillar (R). Aggregate core runtime checks:

$$R = \text{GM}(b_{\text{build}}, b_{\text{runtime}}, b_{\text{type}}, b_{\text{tests}})$$

Step 2: SQL Quality Pillar (S). Weighted combination of SQL metrics:

$$S = 0.50 \times S_1 + 0.30 \times S_2 + 0.20 \times S_4$$

where S_1 (execution correctness) dominates, S_2 (efficiency) contributes secondary value, and S_4 (safety) ensures secure queries.

Step 3: Web Quality Pillar (W). Weighted combination of UI/UX metrics:

$$W = 0.40 \times W_1 + 0.30 \times W_2 + 0.20 \times W_3 + 0.10 \times W_4$$

prioritizing task completion and visual correctness.

Step 4: Agentic DevX Pillar (D).

$$D = \text{GM}(\hat{x}_{\text{run}}, \hat{x}_{\text{deploy}})$$

where \hat{x} = score/5 normalizes 0–5 scores to [0, 1].

Step 5: Soft Penalty Gate. Penalize critical outages without collapsing to zero:

$$G = (0.25 + 0.75 \times b_{\text{build}}) \times (0.25 + 0.75 \times b_{\text{runtime}}) \times (0.50 + 0.50 \times b_{\text{db}}) \times (0.50 + 0.50 \times 1_{S_1 \geq 0.5})$$

Step 6: Final Composite.

$$\text{AppEval-100} = 100 \times (0.30 \times R + 0.25 \times S + 0.25 \times W + 0.20 \times D) \times G$$

Pillar Weight Rationale: Reliability (30%) ensures the app works; SQL Quality (25%) and Web Quality (25%) capture the core value proposition for Databricks data applications; DevX (20%) measures autonomous operability.

Values near 100 denote near-perfect readiness; 50–70 indicates partial operability; <30 signifies fundamental execution issues.

3.5 DORA Metrics Mapping

We map our metrics to industry-standard DORA [3] measures:

- **Deployment Frequency:** Count of successful D9 events per app per evaluation cohort.
- **Lead Time:** Median time from first model call to successful D9 deployment.
- **Change Failure Rate:** Fraction of deployments that fail healthcheck or rollback within 30 min.
- **MTTR:** Median time from failure detection to restore (prior healthy image running).

Production Gate: L1–L7 pass, D8≥4, D9≥4, type-safety pass, and DORA guardrails (Lead Time P50 ≤10 min, CFR ≤15%, MTTR ≤15 min).

3.6 MLflow Integration

We integrate with Databricks Managed MLflow for experiment tracking:

- Automatic metric logging per evaluation run
- Trend analysis across model versions and configurations
- Artifact versioning for reproducibility
- DORA telemetry for delivery performance monitoring

4 Trajectory Optimizer

4.1 Motivation

Agent execution trajectories contain rich signal about tool and scaffold failures that cannot be captured by end-state metrics alone. However, manual analysis of trajectories does not scale.

4.2 Map-Reduce Architecture

We employ a two-phase analysis approach:

Map Phase. Each trajectory is analyzed independently using a fast model (Claude Haiku). The analysis identifies:

- Struggles: errors, retries, confusion patterns
- Friction points: slow progress, repeated attempts
- Inefficient approaches: suboptimal tool usage

Reduce Phase. Individual analyses are synthesized by a more capable model (Claude Opus) with read-only access to the codebase. This phase generates actionable recommendations for:

- Template improvements: structure, guidance, scaffolding
- Tool improvements: missing tools, unclear descriptions
- Root cause analysis: systemic failure patterns

4.3 Feedback Loop

The trajectory optimizer enables a continuous improvement cycle:

Generate → Evaluate → Analyze Trajectories → Improve
Scaffolding → Repeat

5 Experimental Setup

5.1 AppEval-100 Benchmark Design

We design a 100-prompt benchmark targeting statistical validity for Databricks data application evaluation. Our sample size provides ±9% confidence interval at 95% confidence for binary metrics ($p = 0.7$), matching InsightBench [13] (100 tasks) and exceeding Spider 2.0's [8] focused enterprise subset (632 tasks).

Table 2: Prompt difficulty distribution ($n = 100$)

Tier	Count	Description
Simple	40	Single-entity CRUD, basic dashboards, one data source
Medium	40	Multi-entity JOINs, filters, interactive charts, 2–3 data sources
Hard	20	Complex analytics, multi-step workflows, real-time updates

5.1.1 Difficulty Distribution.

465 5.1.2 *Domain Coverage.* Prompts span five application domains:
 466 Analytics Dashboards (26%), CRUD Applications (22%), Data Visualization (22%), Business Intelligence (20%), and Reporting Tools (10%).
 467
 468

469 5.1.3 *Schema Families.* We target six schema families to ensure
 470 diversity: TPC-DS (25 prompts, retail/customer analytics), TPC-H (20,
 471 supply chain), NYC Taxi (15, trip analytics), Custom Databricks (25, Unity Catalog/ML features), Financial (10, trading/risk), and
 472 IoT/Telemetry (5, device metrics).
 473
 474

475 5.2 Prompt Collection Methodology

477 Prompts are collected from three sources to ensure realistic ambiguity
 478 and domain vocabulary:

- 479 (1) **Hackathon recordings** (50 prompts): Extracted from video
 480 transcripts of internal Databricks application hackathons,
 481 preserving natural language vagueness.
- 482 (2) **Production logs** (30 prompts): Anonymized user requests
 483 from app.build, capturing real-world requirements.
- 484 (3) **Synthetic generation** (20 prompts): LLM-generated from
 485 templates to fill coverage gaps.

486 **Quality Criteria:** Prompts must exhibit realistic ambiguity,
 487 domain-specific vocabulary, implicit requirements (unstated but
 488 expected features), and achievability with current templates.
 489

490 5.3 Current Evaluation Dataset

492 For this paper, we evaluate on a “Simple 20” subset—20 Databricks
 493 data application prompts spanning dashboards, analytics, and business
 494 intelligence tools—retained as a fixed regression set for longitudinal
 495 comparison.

497 5.4 Generation Pipeline

498 Applications are generated using:

- 500 • **Claude Agent SDK** with edda MCP for tool orchestration
- 501 • **Dagger** containerized execution for isolation and reproducibility
- 502 • **Environment scaffolding** with templates and validation
 503 pipelines

505 5.5 Statistical Validity

507 Our 100-prompt benchmark provides:

- 508 • **Confidence Interval:** $\pm 9\%$ at 95% confidence for binary
 509 metrics
- 510 • **Statistical Power:** 0.80 for medium effect size ($d = 0.5$)
- 511 • **Stratification:** Three difficulty tiers adequately sampled
 512 (40/40/20)

513 The “Simple 20” regression set enables longitudinal model comparisons while guarding against prompt inflation effects.
 514
 515

516 6 Results

518 6.1 Overall Performance

519 We evaluate on the “Simple 20” prompt set—20 Databricks data
 520 application prompts spanning dashboards, analytics, and business
 521 intelligence tools.

523 **Table 3: Aggregate evaluation results ($n = 20$ applications),
 524 Eval 2.0)**

ID	Metric	Result	Notes	
L1	Build Success	20/20	100% pass	527
L2	Runtime Success	20/20	100% pass	528
L3	Type Safety	1/20	5% pass (improvement needed)	529
L4	Tests Pass	–	Not yet instrumented	530
L5	DB Connectivity	18/20	90% pass	531
L6	Data Operations	–	Requires app-specific procedures	532
L7	UI Validation	–	VLM check in progress	533
D8	Runability	3.0/5	Average score	534
D9	Deployability	2.5/5	Average score	535

536 **Key Finding:** 100% of generated applications achieve build and
 537 runtime success, with 90% achieving functional Databricks connec-
 538 tivity. However, type safety (5%) and agentic DevX scores (3.0/5,
 539 2.5/5) indicate room for improvement toward production readiness.

540 6.2 Generation Efficiency Metrics

541 **Table 4: Efficiency metrics ($n = 20$ applications)**

Metric	Value	Notes	
E10: Total Tokens	16K/app	Prompt + completion	549
E11: Generation Time	6–9 min	End-to-end	550
E12: Agent Turns	93 avg	Conversation turns	551
E13: LOC	732 avg	Lines of code	552
Cost per App	\$0.74	API cost	553
Total Cost (20 apps)	\$14.81	–	554
Build Step Time	2.7s avg	Docker build	555

556 6.3 Comparison: Eval 1.0 vs 2.0

557 **Table 5: Evolution from manual (Eval 1.0) to automated
 558 (Eval 2.0) evaluation**

Aspect	Evals 1.0	Evals 2.0	
Viability Rate	73% (30 apps)	100% build/runtime	566
Time to Deploy	30–60 min	6–9 min	567
Evaluation Method	Manual rubric	Automated pipeline	568
Metrics Tracked	Binary viability	13 metrics + AppEval-100	569
Reproducibility	Low	Full artifact pack	570

571 6.4 Production Readiness Assessment

572 Current status: **below production threshold.** To reach Production
 573 Candidate level:

- 574 • L3 Type Safety: 5% → target $\geq 90\%$
- 575 • D8 Runability: 3.0 → target ≥ 4

- D9 Deployability: $2.5 \rightarrow \text{target } \geq 4$
- DORA guardrails: Lead Time $P50 \leq 10\text{m}$, CFR $\leq 15\%$, MTTR $\leq 15\text{m}$

585 6.5 Trajectory Optimizer Insights

586 Analysis of agent trajectories revealed common friction patterns:

- **SQL Syntax:** Databricks SQL variations causing query failures
- **Error Handling:** Missing error handling in template scaffolding
- **Tool Descriptions:** Unclear MCP tool descriptions leading to incorrect usage
- **Type Inference:** TypeScript strict mode violations in generated code

596 These insights feed back into template and tool improvements via the optimize → evaluate → analyze cycle.

599 7 Discussion

600 7.1 Limitations

602 **Platform Specificity.** Our current implementation targets Databricks applications. Extending to other platforms requires platform-specific metrics (e.g., AWS Lambda, Vercel).

605 **Binary Metrics.** Several metrics are binary, potentially missing nuanced quality differences. Future work could introduce continuous variants.

608 **Dataset Size.** Our evaluation of 20 applications provides initial validation but may not capture edge cases. Scaling to larger datasets is ongoing.

612 7.2 Broader Impact

614 By establishing standardized metrics for autonomous deployability, we enable:

- Reproducible benchmarking of agentic code generation systems
- Objective comparison across different approaches
- Systematic improvement through trajectory-based feedback

621 8 Conclusion

623 We presented Klaudbiusz, an open-source, agent-agnostic toolset for 624 autonomous Databricks application generation. Rather than building a specific agent, we provide reusable infrastructure—scaffolding, 626 templates, and MCP tools—that any agentic system can leverage. Our 628 trajectory analyzer enables continuous improvement of this toolset by identifying friction points in agent execution traces. AppEval-100 provides composite evaluation combining Reliability, 630 SQL Quality, Web Quality, and Agentic DevX, mapped to industry-standard 631 DORA metrics.

632 The path to reliable agentic code generation requires not just better 633 models, but better tooling. We release Klaudbiusz to enable the 634 community to build and improve agent-agnostic infrastructure 635 systematically.

636 **Open Source Release.** Toolset, trajectory analyzer, and evaluation 637 harness available at: [URL redacted for review]

References

- [1] Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. 2021. Program Synthesis with Large Language Models. *arXiv preprint arXiv:2108.07732* (2021).
- [2] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. 2021. Evaluating Large Language Models Trained on Code. *arXiv preprint arXiv:2107.03374* (2021).
- [3] Nicole Forsgren, Jez Humble, and Gene Kim. 2018. *Accelerate: The Science of Lean Software and DevOps: Building and Scaling High Performing Technology Organizations*. IT Revolution Press.
- [4] Xueyu Hu et al. 2024. InfiAgent-DABench: Evaluating Agents on Data Analysis Tasks. *arXiv preprint arXiv:2401.05507* (2024).
- [5] Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. 2024. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? *arXiv preprint arXiv:2310.06770* (2024).
- [6] Jing Yu Koh, Robert Lo, Lawrence Jang, Vikram Duvvur, Ming Lim, Po-Yu Huang, Graham Neubig, Shuyan Zhou, Ruslan Salakhutdinov, and Daniel Fried. 2024. VisualWebArena: Evaluating Multimodal Agents on Realistic Visual Web Tasks. *arXiv preprint arXiv:2401.13649* (2024).
- [7] Yuhang Lai, Chengxi Li, Yiming Wang, Tianyi Zhang, Ruiqi Zhong, Luke Zettlemoyer, Scott Wen-tau Yih, Daniel Fried, Sida Wang, and Tao Yu. 2022. DS-1000: A Natural and Reliable Benchmark for Data Science Code Generation. *arXiv preprint arXiv:2211.11501* (2022).
- [8] Fangyu Lei, Tianbao Xie, et al. 2024. Spider 2.0: Evaluating Language Models on Real-World Enterprise Text-to-SQL Workflows. *arXiv preprint arXiv:2411.07763* (2024).
- [9] Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhu Li, Bowen Li, Bailin Wang, Bowen Qin, Ruixing Geng, Nan Huo, et al. 2023. Can LLM Already Serve as A Database Interface? A Blg Bench for Large-Scale Database Grounded Text-to-SQL. *Advances in Neural Information Processing Systems* 36 (2023).
- [10] Jinyang Li, Nan Huo, Yan Gao, Jiayi Shi, Yingxiu Zhao, Ge Qu, Yurong Wu, Chenhao Ma, Jian-Guang Lou, and Reynold Cheng. 2024. Tapilot-Crossing: Benchmarking and Evolving LLMs Towards Interactive Data Analysis Agents. *arXiv preprint arXiv:2403.05307* (2024).
- [11] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuna Yang, et al. 2023. AgentBench: Evaluating LLMs as Agents. *arXiv preprint arXiv:2308.03688* (2023).
- [12] Grégoire Mialon, Roberto Dessì, Maria Lomeli, Christoforos Nalmpantis, Ram Pasunuru, Roberta Raez, Baptiste Rozière, Timo Schick, Jane Dwivedi-Yu, Asli Celikyilmaz, et al. 2023. GAIA: A Benchmark for General AI Assistants. *arXiv preprint arXiv:2311.12983* (2023).
- [13] Gaurav Sahu et al. 2024. InsightBench: Evaluating Business Analytics Agents Through Multi-Step Insight Generation. *arXiv preprint arXiv:2407.06423* (2024).
- [14] John Yang, Carlos E Jimenez, Alexander Wettig, Kilian Liber, Shunyu Yao, Karthik Narasimhan, and Ofir Press. 2024. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. *arXiv preprint arXiv:2405.15793* (2024).
- [15] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, et al. 2018. Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL Task. *arXiv preprint arXiv:1809.08887* (2018).
- [16] Zhaojian Yu et al. 2024. HumanEval Pro and MBPP Pro: Evaluating Large Language Models on Self-invoking Code Generation. *arXiv preprint arXiv:2412.21199* (2024).
- [17] Shuyan Zhou, Frank F Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Yonatan Bisk, Daniel Fried, Uri Alon, et al. 2024. WebArena: A Realistic Web Environment for Building Autonomous Agents. *arXiv preprint arXiv:2307.13854* (2024).